


Coding exercise: (5 marks) Deadline Apr 30th, 11:59 PM

Take a $n \times n$ grid (Say $n = 5, 10, 20$)



Reward (1 in the top right
0 elsewhere)

π : optimal policy (this easy to compute
just by observation)

* Learn V_π using TD

* choose

$$\phi(s) = (1, x, y, |x-y|, xy, x^2, y^2, x^2+xy^2, (x+y)^2) \\ \in \mathbb{R}^9$$

$$\text{and find } V_\pi(s) \approx \phi^T(s) \theta_*$$

using TD with function approximation.